

Raspberry Pi Cluster Setup Guide

Acknowledgements

Project lead: **Devin Smoot**, *SWOSU, Weatherford, OK*

Contributor/Implementation tester: **Madison Matli**, *SWOSU, Weatherford, OK*

To do list

- Add new sections:
 - Table of contents
 - What this guide covers
 - What you need for this guide (Hardware/software)
 - Who this guide is directed at (audience)
 - Conventions
 - Errata (Include known issues)
 - Fleshed out troubleshooting guide
 - Add NFS section
 - Multi-user setup section
 - Script section for ease of cluster management
 - Look at IPv4 forwarding for better solution
 - Router type setup with computer nodes on LAN side
-

Considerations to Consider before starting

If your SD card size will vary you will want to build the head node using the smallest size of SD card. This will ensure that the image for that SD card will ALWAYS be able to be written to a similar sized SD or larger. If you start with a 64GB SD card you will not be able to write the image to a 16GB SD card.

Head Node

Hardware:

- Raspberry Pi board x 1
- Wi-Fi USB dongle x 1
- SD Card 16GB+ x 1
- Ethernet cable x 1
- HDMI cable x 1
- Power cable mini-USB x 1

Compute nodes

Hardware:

- Raspberry Pi board x 7
- SD Card 16GB+ x 1
- Ethernet cable x 1
- Power cable mini-USB x 1

Additional Hardware

- 10 Port USB hub
- 16 Port gigabit switch

Setup, Installation, and Testing

Step 1 - Install operating systems

Install Raspbian Lite on SD card for head unit(s) and each compute node

[Raspbian Lite](#)

[Raspbian Install Guides](#)

Step 2 - Initial Head Node Setup

Setup the locale settings to make sure the correct keyboard, language, timezone, etc are set. This will ensure we are able to enter the correct symbols while working on the command line.

Configure Locale:

Log in with username: **pi** and password **raspberry**

Start the Raspberry Pi configuration tool:

```
# raspi-config
```

Setup Advanced Options:

Select **7 Advanced Options**

- Select **A3 Memory Split**
 - Enter **16**
 - Press **Enter**

Setup Localisation Options:

Select **4 Localisation Options**

- Select Locale **I1 Change Locale**
 - Unselect **en_GB.UTF-8 UTF-8**
 - Select **en_US ISO-8859-1**
 - Press **Enter**
 - Select **en_US**

Select 4 Localisation Options

- Select **I2 Change Timezone**
 - Select **US** (or appropriate country)
 - Select **Central** (or appropriate local timezone)

Select 4 Localisation Options

- Select **I3 Change Keyboard Layout**
 - Use the default selected Keyboard
 - Press **Enter**
 - Select **Other**
 - Select **English (US)**
 - Select **English (US)**
 - Select **The default for the keyboard layout**
 - Select **No compose key**
 - Press **Enter**

Select 4 Localisation Options

- Select **I4 Change Wi-fi Country**
 - Select **US United States**
 - Select **Ok**

Setup Network Options:

Select 2 Network Options

- Select **N1 Hostname**
 - Select **Ok**
 - Enter **head** for the hostname
 - Press **Enter**

Select 2 Network Options

- Select **N2 Wi-fi**
 - Enter w i-fi SSID
 - Press **Enter**
 - Enter w i-fi passphrase
 - Press **Enter**

Setup Interfacing Options:

Select 5 Interfacing Options

- Select **P2 SSH**
 - Select **Yes**
 - Select **Ok**
 - Press **Enter**

Tab to **Finish**

Select **Yes** to reboot

Step 3 - Configure Network Settings

Edit `/etc/network/interfaces`:

```
# nano /etc/network/interfaces
```

Add the following to the end of the file:

```
auto lo
iface lo inet loopback

iface eth0 inet manual

allow-hotplug wlan0
iface wlan0 inet manual

wpa-conf /etc/wpa_supplicant/wpa_supplicant.conf
```

Setup `eth0` static ip address:

Edit `/etc/dhcpd.conf`:

```
# nano /etc/dhcpd.conf
```

Add to the end of the file:

```
interface eth0
static ip_address=192.168.10.5
static domain_name_servers=8.8.8.8
```

Save and exit

Reboot:

```
# reboot
```

Step 4 - Update the system

```
# apt update && sudo apt upgrade -y
```

Reboot:

```
# reboot
```

Step 5 - IP forwarding for nodes to access internet

Setup IP forwarding so that all compute nodes will have access to the internet for package installation and to download any needed materials on later use.

Log in with username: **pi** and password **raspberry**

Enable IPv4 Forwarding and Disable IPv6:

```
# nano /etc/sysctl.conf
```

Add the following lines to the end of the file (this includes the IP forwarding rule from above):

```
# Enable IPv4 forwarding
net.ipv4.ip_forward = 1

# Disable IPv6
net.ipv6.conf.all.disable_ipv6 = 1
net.ipv6.conf.default.disable_ipv6 = 1
net.ipv6.conf.lo.disable_ipv6 = 1
```

Save and exit

Update the configuration files:

```
# sysctl -p
```

Edit and Save the iptables:

```
# iptables -t nat -A POSTROUTING -o eth0 -j MASQUERADE
# iptables -t nat -A POSTROUTING -o wlan0 -j MASQUERADE

# bash -c "iptables-save > /etc/iptables.rules"
```

Add settings to */etc/network/interfaces* file:

```
# nano /etc/network/interfaces
```

Add the following line at the end of the wlan0 section under wpa-conf line to make the changes persistent:

```
pre-up iptables-restore < /etc/iptables.rules
```

Save and exit

Update */etc/hosts* file:

Add the following to the end of the file:

Note: At this point you want to assign and name all of your nodes that **WILL** be in your cluster and enter them in the hosts file. Below is an example of a 6 node cluster including the head node as one of the six. This file will be copied with the image to the compute nodes and will save you a step of developing and deploying the hosts file later.

Edit `/etc/hosts` file:

```
# nano /etc/hosts
```

Modify or add the following lines to the file:

```
127.0.1.1    head

192.168.10.3    nodeX
192.168.10.5    head
192.168.10.100  node0
192.168.10.101  node1
192.168.10.102  node2
192.168.10.103  node3
192.168.10.104  node4
192.168.10.105  node5
192.168.10.106  node6
```

Reboot:

```
# reboot
```

Install MPICH

Install prerequisite *Fortran* which will be required for compiling MPICH. All other dependencies are already installed.

Step 1 - Install Fortran

```
# apt install gfortran
```

Step 2 - Install and Setup MPICH3

Create hpc group:

```
# groupadd hpc
```

Add pi user to hpc group:

```
# usermod -aG hpc pi
```

Create hpc directory in root:

```
# mkdir -p /software/lib  
$ cd /software/lib
```

Take ownership of /software:

```
# chown -R pi:hpc /software
```

Create build and install directory inside mpich3 directory:

```
$ mkdir mpich_3  
$ cd mpich_3  
$ mkdir build install
```

Download mpich3 and untar:

```
$ wget http://www.mpich.org/static/downloads/3.3/mpich-3.3.tar.gz  
$ tar xvfz mpich-3.3.tar.gz
```

Compile and install mpich3:

```
$ cd build  
$ /software/lib/mpich_3/mpich-3.3/configure --prefix=/software/lib/mpich_3/install  
$ make  
$ make install
```

Activate environment variable:

```
$ export PATH=/software/lib/mpich_3/install/bin:$PATH
```

Add path to environment variables for persistence:

```
# nano ~/.bashrc
```

Add the following to the end of the file:

```
# MPICH
export PATH="/software/lib/mpich_3/install/bin:$PATH"
```

Step 3 - Create list of nodes for MPI:

This list of nodes will need to be updated as you add nodes later. Initially you will only have the head node.

Create node list:

```
$ cd ~
# nano nodelist
```

Add the head node ip address to the list:

```
192.168.10.5
```

Note: Anytime you need to add a node to the cluster make sure to add it here as well as */etc/hosts* file.

Step 4 - Test MPI

Test 1 - Hostname Test

Enter on command line:

```
$ cd ~
$ mpiexec -f nodelist hostname
```

Output:

```
head
```

Test 2 - Calculate Pi

Enter on command line:

```
$ mpiexec -f nodelist -n 2 /software/lib/mpich_3/build/examples/cpi
```

Output:

```
Process 0 of 2 is on head
Process 1 of 2 is on head
pi is approximately 3.1415926544231318, Error is 0.000000008333387
wall clock time = 0.003250
```


Step 5 - Setup SSH keys

Note: Must be executed from head node as pi user

Generate SSH key:

```
$ cd ~  
  
$ ssh-keygen -t rsa -P "" -f ~/.ssh/id_rsa -q
```

Transfer the key to the authorized_keys file:

```
$ cat ~/.ssh/id_rsa.pub > ~/.ssh/authorized_keys
```

Prepare for cloning

Shutdown the head node:

```
# shutdown -h now
```

Save SD Image

At this point you will want to save an image of the head node. This will give you a fall back point if you make mistakes moving forward. You will also use this image to begin your node image.

Using the same guide as described in the beginning you will want to reverse the process of writing an image to the SD and *read* an image from the SD and save that image to your PC. Now you have saved your SD like a checkpoint.

Sample name for SD image:

```
compute_node_mpi_stage_2017_01_03
```

Create Node image

The overview of this process:

1. Save image of *head node*.
2. On a new SD card write the *head node* image you just saved.
3. Boot the second SD you just created from the head node and make the following changes for "Creating a Generic Node Image".
4. Save image of newly created *generic compute node*.

At this point you have a copy of both the *head node* and *generic compute node* at the MPI stage. This is a checkpoint that you can fall back to if there are errors after this point.

Create Generic Node image

Completing this step will give you a node image that can be quickly written to an SD card and distributed to expand your cluster. This will be a repeatable process when completed. You will setup an initial *compute node* image using your saved *head node* image. You will go in and change the hostname, hosts file (to match the hostname), and ip address (to a generic, always easy to find ip address for when you need to configure the node after deployment) to *generic settings*. Doing this will allow you to always access your *generic compute node* image at the same IP address and hostname. You will then be able to set up the compute node image to a specific IP address and hostname. Following this process will allow for prompt and efficient deployment of a cluster.

Step 1 - Boot image and login

Log in with username: **pi** and password **raspberrypi**

Step 2 - Enter a generic ip address

```
# nano /etc/dhcpd.conf
```

Change the *eth0* ip address from:

```
static ip_address=192.168.10.5
```

To:

```
static ip_address=192.168.10.3
```

Also add to the end of the file:

```
static routers=192.168.10.5
```

Save and exit

Step 3 - Enter a generic hostname

```
# nano /etc/hostname
```

Change:

```
head
```

To:

```
nodeX
```

Save and exit

Step 4 - Edit hosts file

```
# nano /etc/hosts
```

Change:

```
127.0.1.1          head
```

To:

```
127.0.1.1          nodeX
```

Save and exit

Step 5 - Remove wireless connection information

Edit *interfaces* file:

```
# nano /etc/network/interfaces
```

Remove:

```
allow-hotplug wlan0
iface wlan0 inet manual

wpa-conf /etc/wpa_supplicant/wpa_supplicant.conf

pre-up iptables-restore < /etc/iptables.rules
```

Edit *wpa_supplicant.conf*:

```
# nano /etc/wpa_supplicant/wpa_supplicant.conf
```

If you are using a secure network remove this section:

```
network={
    ssid="<network name>"
```

```
psk="<network password>"
}
```

If you are using an unsecure network remove this section:

```
network={
  ssid="<network name>"
  key_mgmt=NONE
}
```

Step 6 - Shutdown and create a new image of the SD

```
# shutdown -h now
```

Now you will go back to WinDiskImager32 and save the image as a node image. This is a generic node image that you can quickly deploy and use to set up your cluster with.

Sample name for SD image:

```
compute_node_mpi_stage_2017_01_03
```

Setup Generic Node image

In this step you will configure a newly deployed node from your generic node image you created. With this you will ssh in to the generic node's ip address and configure the hostname, hosts file, and ip address. Each of these will be set to the node's new permanent settings within the cluster. When completing this step if you are deploying multiple new nodes you will need to power them up one at a time and configure them one at a time. This is due to all of the nodes using the same ip address. This allows for all nodes to use a single point of reference to quickly find and deploy them.

[Raspbian Install Guides](#)

Step 1 - Copy generic node image created earlier to an SD card using WinDiskImager32.

Step 2 - Boot and login to your system

Log in with username: **pi** and password **raspberrypi**

SSH into the new node:

```
$ ssh pi@nodeX
```

Enter **yes** to accept the key

Verify you are logged in:

Command prompt should read `pi@nodeX:~ $`

Step 3 - Adjust */etc/hostname* file

```
# nano /etc/hostname
```

Change:

```
nodeX
```

To:

```
node0
```

Save and exit

Note: This number will increment by one each time you add a node and must be unique on your cluster.

Step 4 - Adjust */etc/dhcpd.conf*

```
# nano /etc/dhcpd.conf
```

Change the *eth0* ip address from:

```
static ip_address=192.168.10.3
```

To:

```
static ip_address=192.168.10.100
```

Save and exit

Step 5 - Edit hosts file

```
# nano /etc/hosts
```

Change:

```
127.0.1.1          nodeX
```

To:

127.0.1.1

node0

Save and exit

Step 6 - Expand Filesystem

Open configuration tool:

```
# raspi-config
```

Select **7 Advanced Options**

Select **A1 Expand Filesystem**

Select **Ok**

Tab to Select **Finish**

Select **Yes**

All settings should take effect on reboot

Deploy Head Node SSH Key

Issue the following command from the head node for each node in the cluster:

Only run this command once the node is restarted with a node number.

```
# rsync -a --rsync-path="sudo rsync" ~/.ssh/authorized_keys pi@node0:~/.ssh/authorized_keys
```

Note: At this point you will just do this once to develop a compute node image with Slurm installed. After that is complete you will create a new generic image of the compute node. Once that is complete you can use that image to finish deploying your compute nodes for the rest of your cluster.

SSH in to the new node:

```
$ ssh pi@nodeX
```

Reboot the node:

```
# reboot
```

Install NTP

NTP is used to keep the cluster time close together using outside NTP servers to sync with the head node. All computer nodes will sync with the head node.

Reference:

<http://raspberrypi.tomasgreno.cz/ntp-client-and-server.html>

<http://www.pool.ntp.org/zone/north-america>

Head Node

Install NTP:

```
# apt install ntp
```

Edit the */etc/ntp.conf*:

```
# nano /etc/ntp.conf
```

Change:

```
pool 0.debian.pool.ntp.org iburst
pool 1.debian.pool.ntp.org iburst
pool 2.debian.pool.ntp.org iburst
pool 3.debian.pool.ntp.org iburst
```

To:

```
server 0.north-america.pool.ntp.org
server 1.north-america.pool.ntp.org
server 2.north-america.pool.ntp.org
server 3.north-america.pool.ntp.org
```

Restart NTP:

```
# /etc/init.d/ntp restart
```

Compute Node

SSH in to compute node:

```
$ ssh pi@node0
```

Install NTP on comput node:

```
# apt install ntp
```

Set Head Node as NTP server.

Edit */etc/ntp.conf*:

```
# nano /etc/ntp.conf
```

Under `restrict ::1` add:

```
restrict 192.168.10.0 mask 255.255.255.0
```

Change:

```
#broadcast 192.168.123.255
```

To:

```
broadcast 192.168.10.255
```

Save and exit

Restart NTP service:

```
# /etc/init.d/ntp restart
```

Exit to head node:

```
$ exit
```

Install Slurm on Head Node

Slurm is the scheduler that organizes jobs to be run on the cluster. This interfaces with MPI and finds the most efficient ways to run jobs according to available resources. This must be installed after creating the base generic image as there is a Slurm controller that must be run on the head node. This is different from the Slurm client that is run on compute nodes.

[Slurm scontrol command]<https://slurm.schedmd.com/scontrol.html>

[Slurm configuration information]https://wiki.fysik.dtu.dk/niflheim/Slurm_configuration

Step 1 - Install Slurm


```
# apt install slurm-wlm slurmctld
```

Step 2 - Add configuration file

Create the new Slurm configuration file */etc/slurm-llnl/slurm.conf*:

```
# nano /etc/slurm-llnl/slurm.conf
```

Add the following to the file and save:

```
# slurm.conf file generated by configurator easy.html.
# Put this file on all nodes of your cluster.
# See the slurm.conf man page for more information.
#
ControlMachine=head
ControlAddr=192.168.10.5
#
#MailProg=/bin/mail
MpiDefault=none
#MpiParams=ports=-#
ProctrackType=proctrack/pgid
ReturnToService=2
SlurmctldPidFile=/var/run/slurm-llnl/slurmctld.pid
#SlurmctldPort=6817
SlurmdPidFile=/var/run/slurm-llnl/slurmd.pid
#SlurmdPort=6818
SlurmdSpoolDir=/var/lib/slurm/slurmd
SlurmUser=slurm
#SlurmdUser=root
StateSaveLocation=/var/lib/slurm/slurmctld
SwitchType=switch/none
TaskPlugin=task/none
#
#
# TIMERS
#KillWait=30
#MinJobAge=300
#SlurmctldTimeout=120
#SlurmdTimeout=300
#
#
# SCHEDULING
FastSchedule=1
SchedulerType=sched/backfill
#SchedulerPort=7321
SelectType=select/linear
#
#
# LOGGING AND ACCOUNTING
AccountingStorageType=accounting_storage/none
ClusterName=raspi3
#JobAcctGatherFrequency=30
JobAcctGatherType=jobacct_gather/none
#SlurmctldDebug=3
SlurmctldLogFile=/var/log/slurm/slurmctld.log
#SlurmdDebug=3
SlurmdLogFile=/var/log/slurm/slurmd.log
#
#
# COMPUTE NODES
NodeName=node[0-6] Procs=1 RealMemory=768 State=UNKNOWN
```

```
PartitionName=raspi3 Default=YES Nodes=node[0-6] State=UP MaxTime=INFINITE
```

Check if Slurm controller is running:

```
# scontrol show daemons
```

Output:

```
slurmctld
```

Step 3 - Create Munge key

```
# /usr/sbin/create-munge-key
```

Agree to overw rite.

Step 4 - Create log folder and take ownership

```
# mkdir /var/log/slurm
# chown -R slurm:slurm /var/log/slurm/
```

Step 5 - Finish installs and start services

```
# systemctl enable slurmctld.service
# ln -s /var/lib/slurm-llnl /var/lib/slurm
# systemctl start slurmctld.service
# systemctl enable munge.service
# systemctl restart munge.service
```

Verify Slurm controller is running:

```
# systemctl status slurmctld.service
```

Will return feedback to the screen. Verify *Active* line states: **active (running)**.

Verify Munge is running:

```
# systemctl status munge.service
```

Will return feedback to the screen. Verify *Active* line states: **active (running)**.

Step 6 - Add user to Slurm group

```
# adduser pi slurm
```

Check Slurm status:

```
$ sinfo
```

Output:

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
raspi3*	up	infinite	7	idle	node[0-6]

Install Slurm on Compute Node

Step 1 - Copy Slurm configuration and Munge files from *Head Node*

On *head node*:

Give rsync proper permission to run:

```
# visudo
```

Add the following to the end of the file:

```
<username> ALL=NOPASSWD: /usr/bin/rsync *
```

On *compute node*:

SSH in to compute node:

```
$ ssh pi@node0
```

Give rsync proper permission to run:

```
# visudo
```

Add the following to the end of the file:

```
<username> ALL=NOPASSWD: /usr/bin/rsync *
```

Exit back to head node:

```
$ exit
```

Step 2 - Install Slurm daemon

Execute on *node0*:

SSH in to *node0*:

```
$ ssh pi@node0
```

Install Slurm daemon and Slurm client:

```
# apt install slurmd slurm-client  
# ln -s /var/lib/slurm-llnl /var/lib/slurm
```

Create log folders and take ownership:

```
# mkdir -p /var/log/slurm  
# chown -R slurm:slurm /var/log/slurm
```

Take ownership of Slurm run folder:

```
# chown -R slurm:slurm /var/run/slurm-llnl
```

Exit to head node:

```
$ exit
```

On *head node*:

Copy Munge and Slurm configuration files from head node to compute node:

See "Generate and distribute Munge key script" in the Troubleshooting section

```
# rsync -a --rsync-path="sudo rsync" /etc/slurm-llnl/slurm.conf pi@node0:/etc/slurm-llnl/slurm.conf
```

On compute node:

SSH in to *node0*:

```
$ ssh pi@node0
```

Take ownership of *munge.key* file:

```
# chown munge:munge /etc/munge/munge.key
```

Finish install and start Slurm and Munge:

```
# systemctl enable slurmd.service
# systemctl restart slurmd.service
# systemctl enable munge.service
# systemctl restart munge.service
```

Verify Slurm daemon is running:

```
# systemctl status slurmd.service
```

Will return feedback to the screen. Verify *Active* line states: **active (running)**.

Verify Munge is running:

```
# systemctl status munge.service
```

Will return feedback to the screen. Verify *Active* line states: **active (running)**.

Step 3 - Add user to Slurm group

```
# adduser pi slurm
```

Execute on *head node*:

```
# scontrol reconfigure
```

Check node status:

```
$ sinfo
```

This should show all nodes in an idle state.

Deploying the Rest of the Cluster

By now you have developed a head node image that contains both MPI and Slurm. You have also developed a compute node image that contains both MPI and Slurm as well. Now you should go back to the instructions for "Create Node Image" to save both images and then use the compute node image to finish deploying your cluster. Saving these images at each stage gives you different configurations that you can easily deploy in the future and also allows you to have a checkpoint in case something goes wrong. You can write the saved node image to your SD and start from that point rather than starting from the beginning.

You should now have a working Raspberry Pi cluster.

Add an ethernet adapter

Add *eth0*:

Edit */etc/network/interfaces* file:

```
# nano /etc/network/interfaces
```

Add below *eth0* section:

```
auto eth1
iface eth1 inet manual
```

Change or add iptables rule to end of file:

```
pre-up iptables-restore < /etc/iptables_wired.rules
```

Edit */etc/dhcpd.conf* file:

Add the following to the end of the file:

```
interface eth1
static ip_address=192.168.1.XXX
static domain_name_servers=8.8.8.8
static routers=192.168.1.1
```

Create iptables rules file:

Flush the iptables in Memory

```
sudo iptables -F
```

Rebuild the rules and file:

```
sudo iptables -t nat -A POSTROUTING -o eth0 -j MASQUERADE
sudo iptables -t nat -A POSTROUTING -o eth1 -j MASQUERADE
```

Save the iptables_wired.rules file:

```
sudo bash -c "iptables-save > /etc/iptables_wired.rules"
```

Check for the iptable rules in the */etc/network/interfaces* file:

Make sure that the line below is present and not commented out:

```
pre-up iptables-restore < /etc/iptables_wired.rules
```

Disable *wlan0*:

Edit */etc/wpa_supplicant/wpa_supplicant.conf* file:

```
# nano /etc/wpa_supplicant/wpa_supplicant.conf
```

Comment out the `network={ connection information }` section (all lines)

Now all traffic for the cluster is routed through eth0 and out eth1 to the internet. Any returning traffic or downloads come in via eth1 and through eth0 to the cluster unless its meant for the head node.

NFS

Choose a disk:

You can use either a standard external hard drive, or a USB flash drive. If using an external hard drive you will want one that has its own power plug. Drawing power from the Raspberry Pi may cause an undervoltage situation.

Formatting the disk:

For this use we will be formatting in FAT32 for ease of use.

Plug in the drive to a USB port.

Install the software needed and format the USB drive:

```
# apt-get install dosfstools
# mkfs.vfat /dev/sda1 -n USB
```

Mount the disk to */users* directory:

```
# mkdir /users
# chown -R pi:hpc /users
# mount /dev/sda1 /users -o uid=pi,gid=pi
```

Add automatic mounting on boot:

Add the following to */etc/fstab* file:

```
/dev/sda1 /users auto defaults,user 0 1
```

Install the NFS server on the head node:

```
# apt install nfs-server
```

Add the following to the */etc/exports* file:

```
/users 192.168.10.5/24(rw,sync)
```

Restart RPC services:

```
# update-rc.d rpcbind enable && sudo update-rc.d nfs-common enable
# reboot
```

Mount the disk from another Raspberry Pi node:

Install required software:

```
# apt install nfs-common autofs
```

Create the mount point:

```
# mkdir /users
# chown -R pi:hpc /users
```

Add the following to the */etc/auto.master* file:

```
/mnt/nfs /etc/auto.nfs
```

Create the */etc/auto.nfs* file and add the following:

```
pi 192.168.10.5:/users
```


Restart the `autofs` service:

```
# /etc/init.d/autofs restart
```

```
192.168.10.5:/users  /users  nfs      auto    0      0
```

<https://raspberrypi.stackexchange.com/questions/87057/cannot-automatically-mount-nfs-share-to-raspberry-pi>

```
sudo mkdir /software/scripts
sudo mkdir /software/files
```

```
sudo nano ~/.bashrc
```

```
# SCRIPTS
export PATH="/software/scripts:$PATH"
```

This script will deploy files to all nodes to a folder defined by the user. This is a workaround for a multi-user environment.

```
#!/bin/bash

if [ "$1" == "-help" ] || [ "$1" == "" ]; then
    echo -e "Command      \tExample"
    echo "-----"
    echo "deploy_file  deploy_file <filename> <destination folder>"
    echo -e "\t\tNote:Default destination folder is '/software/files'"
    echo "Help          deploy -help"
    exit
fi
```

```

if [ "$1" != "" ]; then
  if [ "$2" == "" ]; then
    filelocation=/software/files
  else
    filelocation=$2
  fi
  echo "Transferring file: $1 to node0:$filelocation"
  rsync $1 pi@node0:$filelocation
  echo "Transferring file: $1 to node1:$filelocation"
  rsync $1 pi@node1:$filelocation
  echo "Transferring file: $1 to node2:$filelocation"
  rsync $1 pi@node2:$filelocation
  echo "Transferring file: $1 to node3:$filelocation"
  rsync $1 pi@node3:$filelocation
  echo "Transferring file: $1 to node4:$filelocation"
  rsync $1 pi@node4:$filelocation
  echo "Transferring file: $1 to node5:$filelocation"
  rsync $1 pi@node5:$filelocation
  echo "Transferring file: $1 to node6:$filelocation"
  rsync $1 pi@node6:$filelocation
else
  echo "All nodes are already defined in the script"
  echo "Please enter a filename and destination folder: ie. deploy_file <filename> <destination folder>"
fi

```

Mutli-user setup script

This will create and configure users from the head node. It will also remove users with options.

```

#!/bin/bash

# Creation script for user creation on SWOSU Raspberry Pi 3 cluster

# Usage display

# Local user creation

# Remote user creation on all nodes

# Set required environment variables

# Help display

```

Troubleshooting Section

Generate and distribute new SSH key script

Generate and distribute new Munge key script

Create a new file in pi user home directory called munge-key-gen.sh:

```

nano ~/munge-key-gen.sh

#!/bin/bash

# Create a new munge key on head node
sudo /usr/sbin/create-munge-key

```

```
for i in {0..6}
do
    # Copy to compute nodes
    sudo cat /etc/munge/munge.key | ssh pi@node$i "sudo cat >> ~/munge.key"

    # Remove old munge.key file
    ssh pi@node$i "sudo rm /etc/munge/munge.key"

    # Move from /home/pi/ to /etc/munge/
    ssh pi@node$i "sudo mv ~/munge.key /etc/munge/munge.key"

    # Take ownership of munge.key by munge user
    ssh pi@node$i "sudo chown munge /etc/munge/munge.key"

    # Change permissions of group and world
    ssh pi@node$i "sudo chmod g-rw,o-rw /etc/munge/munge.key"

    # Restart munge service
    ssh pi@node$i "sudo service munge restart"
done
```

Received SIGHUP or SIGTERM from Nano

Enter the command:

```
bash
```

NETWORK UNREACHABLE:

When experiencing network connectivity problems with compute nodes:

1. Flush the iptables in Memory

```
sudo iptables -F
```

2. Rebuild the rules and file

Repeat the IP tables section of the guide, starting with the commands:

```
sudo iptables -t nat -A POSTROUTING -o eth0 -j MASQUERADE
sudo iptables -t nat -A POSTROUTING -o eth1 -j MASQUERADE
```

3. Save the iptables_wired.rules file:

```
sudo bash -c "iptables-save > /etc/iptables_wired.rules"
```

4. Check for the iptable rules in the */etc/network/interfaces* file:

Make sure that the line below is present and not commented out:

```
pre-up iptables-restore < /etc/iptables_wired.rules
```

5. Reboot:

```
sudo reboot
```

If it is missing then add it to the end of the file. Save and exit.

RSYNC ISSUES:

If having trouble with using rsync commands:

Setup Rsync:

On *Both nodes*

Edit the `/etc/sudoers` file:

```
sudo visudo
```

Add this line to the end of the file:

```
<username> ALL=NOPASSWD: /usr/bin/rsync *
```

MPI ISSUES

If `mpiexec` command fails to execute, stalls, or displays an error message about an unreadable path file:

- `Mpich3` could be in the wrong directory
- Make sure the export path correlates to the actual install path for `MPICH3`
- Reinstalling `MPICH3` and setting up the proper environment variables can fix many problems, re-evaluate the `MPICH3` install instructions and verify all settings before attempting a reinstall.

SSH ISSUES

If the PI is displaying SSH errors when running the `mpiexec` command:

Check the problematic node's `authorized_keys` file, and compare it with the head node's `authorized_keys` file.

Check the file by going to the SSH directory:

```
cd ~/.ssh
```

Now check the file information for `authorized_keys` file:

```
ls -ls
```

The filesize is listed after the owner and group names.

These files should be identical in length, if not redistribute the head node's `authorized_keys` file to the compute node using the following command:

```
rsync -a --rsync-path="sudo rsync" ~/.ssh/authorized_keys pi@nodeX:~/.ssh/authorized_keys
```

COMMANDS TO CHECK SERVICE STATUSES

These commands do the same thing, just with a different syntax:

```
sudo systemctl [start,stop,restart,status] <service name>

sudo service <service name> [start,stop,restart,status]

sudo /etc/init.d/<service name> [start,stop,restart,status]
```

ENABLING/DISABLING NETWORK INTERFACE CONNECTIONS

This is a quick way to bring down and bring back up network interfaces without restarting.

Disable the specified connection

```
sudo ifdown <connection name>
```

Enable the specified connection

```
sudo ifup <connection name>
```

SLURM ISSUES

Make sure the `slurm.conf` file is identical across all nodes.

Use `sudo scontrol reconfigure` to distribute the `slurm.conf` from the head node to all compute nodes responding.

When running the service status command, read the error messages that are displayed: ***these messages are vital in order to troubleshoot current problems.***

PROBLEMATIC NODES

On many occasions, certain nodes fail to work because of a software/hardware malfunction. This can be fixed by removing and reinstalling the software. Hardware problems can be fixed by reformatting the node's SD card, and rewriting it with a functional node image. Also check each Ethernet cable for weaknesses, and verify that each node in the cluster is properly connected.

-For Pi 3 Clusters: The head node is connected via Wi-Fi, and each compute node uses the head node's wireless connection to download files.

-For Pi 2 Clusters: A Wi-Fi adapter is a tested solution for establishing a wireless connection with a Raspberry Pi model 2. Using other wireless adapters could result in incompatible drivers or other various issues. The head node can also be connected to the Internet via an Ethernet cable.

Network Diagrams


Base Equipment Layer (Pictured Below)




Physical Layer (Pictured Below)



Logical Layer (Pictured Below)



Physical and Logical Layers (Pictured Below)



References

<https://www.modmypi.com/blog/how-to-give-your-raspberry-pi-a-static-ip-address-update>

<https://www.raspberrypi.org/forums/viewtopic.php?f=28&t=44609>

<http://www.orkwplay.com/play/automatically-connect-a-raspberry-pi-to-a-wifi-network/>

<https://www.raspberrypi.org/forums/viewtopic.php?f=36&t=162096>

<https://www.raspberrypi.org/forums/viewtopic.php?t=118804&p=808453>

<https://www.raspberrypi.org/forums/viewtopic.php?f=28&t=37575>

<http://unix.stackexchange.com/questions/88100/importing-data-from-a-text-file-to-a-bash-script>

<http://www.tldp.org/LDP/abs/html/arrays.html>

http://how-towiki.com/wiki/How_to_read_command_line_arguments_in_a_bash_script

<http://ryanstutorials.net/bash-scripting-tutorial/bash-variables.php>

<http://stackoverflow.com/questions/19996089/use-ssh-to-start-a-background-process-on-a-remote-server-and-exit-session>

<http://stackoverflow.com/questions/29142/getting-ssh-to-execute-a-command-in-the-background-on-target-machine>

<http://www.igeekstudio.com/blog/setting-up-ganglia-and-hadoop-on-raspberry-pi>

<http://ccm.net/faq/2540-linux-create-your-own-command>

<https://github.com/XavierBerger/RPi-Monitor>

<https://www.raspberrypi.org/blog/visualising-core-load-on-the-pi-2/>

<https://github.com/davidsblog/rCPU>

<https://raseshmori.wordpress.com/2012/10/14/install-hadoop-nextgen-yarn-multi-node-cluster/>

<https://www.packtpub.com/hardware-and-creative/raspberry-pi-super-cluster>